

CHRISTOPHER WANG

1-613-986-4968 | christopher.jj.wang@gmail.com | christopher-wang.github.io

EDUCATION

Software Engineering - Bachelor of Engineering | Carleton University | Ottawa, Ontario 09/2016 – 08/2021

- **Cumulative GPA:** 11.88/12 (A+) | [Chancellor's Medal](#) (rank 2 of 795 undergraduate students) | Dean's List
- **Student Involvement:** Carleton AI Society (President), Carleton Kaggle Club (President)

RESEARCH EXPERIENCE

Undergraduate Student Researcher | Carleton University | Ottawa, Ontario (Part Time) 05/2018 – 04/2021

- Applied unsupervised image segmentation by backpropagation models to multi-channel satellite imagery
- Conducted hierarchical remote sensing on environmental satellite imagery using unsupervised random forests as an embedding algorithm on pixel-wise channel features and agglomerative clustering, achieving an accuracy of ~70%

Research Scientist Intern | Nuance Communications | Montreal, Quebec (Remote) 05/2020 – 08/2020

- Implemented homomorphically encrypted representation search using Face-Net embeddings and Microsoft SEAL
- Enabled an efficient homomorphically encrypted dot product by integrating the vector sum operation into the Python bindings of the TenSEAL open-source library, facilitating nearest neighbor search on embedding tensors
- Experimented with polynomial activation functions for fully encrypted convolutional neural network inference
- Performed literature review on private machine learning inference with a focus on cryptographic solutions

Undergraduate Student Researcher | Carleton University | Ottawa, Ontario (Part Time) 09/2019 – 04/2020

- Fine-tuned pretrained Resnet-18 models to orca whale call spectrograms supplied by Fisheries and Oceans Canada and the INTERSPEECH 2019 ComParE Orca Activity Subchallenge achieving an AUC ROC of 0.89
- Developed a gradient-boosted tree model trained on short-window audio features achieving an AUR ROC of 0.72

Undergraduate Student Researcher | Carleton University | Ottawa, Ontario (Full Time) 05/2018 – 12/2018

- Created an ensemble of U-Net Architectures of varying feature extractors (inception modules, residual blocks, dense blocks, attention connections, etc) for left atrial MRI segmentation attaining a Dice score of 0.92%
- Developed a framework for rapid prototype testing of image segmentation models with data augmentation (translation, rotation, gaussian blur, elastic transform) for 2D and 3D MRI images built on Keras/TensorFlow

PAPERS & PATENTS

Systems and methods for dynamic demand sensing and forecast adjustment (US20230401592A1)	08/2023
Systems and methods for dynamic demand sensing (US20210110413A1)	11/2022
Reciprocal perspective as a super learner improves drug-target interaction prediction (Scientific Reports)	08/2022
Canadian Jobs amid a Pandemic (IEEE COMPSAC)	07/2022
Systems and methods for features engineering (US20230086226A1)	04/2020
U-Net Architecture Ensembles for Left Atrial Segmentation (SPIE Medical Imaging)	06/2019

AWARDS, FUNDING, & SCHOLARSHIPS

Chancellor's Medal (Awarded to the undergraduate student with the second-highest overall average)	10/2021
Internship-Carleton University Research Experience for Undergraduate Students (\$2250)	09/2020 – 04/2021
Internship-Carleton University Research Experience for Undergraduate Students (\$2250)	09/2019 – 04/2020
Internship-Carleton University Research Experience for Undergraduate Students (\$2250)	09/2018 – 04/2019
NSERC Undergraduate Student Research Award (\$6300)	05/2018 – 08/2018
General Dynamics Scholarship (\$5000)	09/2016

WORK EXPERIENCE

Machine Learning Engineer | Kinaxis | Ottawa, Ontario (Full Time) 05/2021 - Present

- Designed and implemented, a retrieval-augmented generative AI backend system using the OpenAI API, FastAPI, Azure Cognitive Search, and MongoDB, with automatic CI/CD integration using Docker, Helm, and Github Actions
- Crafted a novel New Product Introduction (NPI) algorithm for demand forecasting using the PySpark DataFrame API, increasing the accuracy for new product forecasts by 35% (from 0.85 WMAPE to 0.50 WMAPE)

- Architected a data segmentation module using the PySpark, and Argo Workflows, allowing for the processing of terabyte-scale data, resulting in a 8.19x end-to-end speed up (from 131 hours to 16 hours), 18.84x reduction in data storage (from 6.8 TB to 361GB), and a 10.11x reduction in yearly cloud costs (from 1,377,012\$ to 136,084\$)
- Led end-to-end ownership of the forecast explainability feature based on SHapely Additive exPlanations (SHAP)
- Supervised an intern project for accelerating gradient boosted tree training using self adaptive learning rates

Machine Learning Developer Co-Op | Kinaxis | Ottawa, Ontario (Full Time) 01/2019 – 12/2019

- Prototyped various architecture designs for an automatic time series forecasting using technologies like Apache Kafka, Spark, Dask, Pandas, LightGBM, PrestoDB, Amazon S3, and Postgresql/PostGIS on a Kubernetes Cluster
- Experimented with neural networks (LSTMs, 1D convolutions), ensemble tree models, SVM regressors and feature selection techniques (genetic algorithms, mutual information, Boruta) for times series forecasting
- Introduced Dask on Kubernetes for parallelizing python jobs, resulting in linear speed ups (up to 16x)

SELECTED PROJECTS

Projection Efficient Intrinsic Dimensionality | Personal Research 05/2023

- Experimented with LOw Rank Adaptation (LORA) style down projections prior to learning with intrinsic dimensionality, reducing the dense projection size by over 15x with negligible impact to accuracy on MNIST dataset

Implicit Neural Neural Representation (INNR) Models | Personal Research 01/2023

- Implemented an implicit neural representation (INR) model to encode the parameters of a LeNet5 variant trained on MNIST dataset, reducing parameter size by over 35% with minimal accuracy degradation (95.56% to 95.00%)

Inebriation Prediction using Keystroke Dynamics | Personal Research 10/2021

- Engineered a web application using Vue.js and Flask to collect keystroke data of inebriated vs sober users
- Applied a feedforward neural network to classify inebriation from keystroke dynamic features such as hold time, press-press time, release-release time, release-press time and press-release time, achieving an accuracy of 95.83%

Transformers for Protein Function Prediction | Capstone 04/2021

- Formulated protein ontology prediction as an eXtreme MultiLabel Classification problem, applied probabilistic label trees and fine-tuned ProtBERT models beating the DeepGoPlus model baseline across all ontology domains

SHapely Additive exPlanations for Metric Explainability | Industry Research 04/2020

- Developed a novel formulation of Kernel SHapely Additive exPlanations for probabilistic and bayesian model explanations using a Kullback–Leibler divergence metric, submitted to CCAI 2020 but was weakly rejected

CANSSI National Case Study Competition | National Student Competition 09/2019

- Achieved Silver at the national level poster competition, presented at Simon Fraser University
- Achieved Gold at the provincials level Kaggle competition for real-time ferry lateness prediction by applying bayesian hyperparameter optimization, feature engineering, HSIC Lasso feature selection, and tree ensembles

Jigsaw Unintended Bias in Toxicity Classification | Kaggle Competition 06/2019

- Achieved Bronze (top 10% of competitors, placing 195th of 3165 competitors) for toxic comment classification with BERT fine tuning, model ensembling, dynamic mini-batching, and proportioned head/tail text truncation

DistriFlow | Software Project 02/2019

- Devised a Node.js module that implements federated learning and asynchronous stochastic gradient descent for distributing neural network training on web browsers, built using Typescript, socket.io and TensorFlow.js

TECHNICAL SKILLS

- **Programming Languages:** Python, Java, SQL, R
- **Analytics/ML Libraries:** Pandas, Pytorch, PySpark, SciKit-Learn, LightGBM, HuggingFace, Keras, TensorFlow
- **Deployment Technologies:** Docker, Kubernetes, Helm, Azure, GCP, Databricks, Argo, GitHub Actions, SLURM